

```
import os
import shutil
import pandas as pd
import numpy as np
import requests
from sklearn.model_selection import train_test_split
from lightgbm import LGBMRegressor
import shap
import matplotlib.pyplot as plt # For visualizations
import seaborn as sns

supress_warnings = True

# Set seed for reproducibility
np.random.seed(42)
```

▼ Data Loading and Preparation

```
import pandas as pd
import numpy as np
import requests
from io import BytesIO
import zipfile
import os

def download_csv(url, filename):
    print(f"Downloading {filename} ...")
    r = requests.get(url)
    r.raise_for_status()
    with open(filename, "wb") as f:
        f.write(r.content)
    print(f"Saved to {filename}")

places_url = "https://data.cdc.gov/api/views/cwsq-ngmh/rows.csv?access_type=CSV&method=GET"
places_filepath = "/Users/sankates/workspace/MSAI/AI-391MCSML/PLACES.csv"

# Add check for file existence
if not os.path.exists(places_filepath):
    download_csv(places_url, places_filepath)
else:
    print(f"File already exists at {places_filepath}. Skipping download.")

places = pd.read_csv(places_filepath, dtype={'LocationID':'str'})
places = places.rename(columns={'LocationID': 'fips'})
```

File already exists at /Users/sankates/workspace/MSAI/AI-391MCSML/PLACES.csv

```
# Drop unnecessary columns
places = places.drop(columns=['Year', 'DataSource', 'StateDesc', 'Data'])

# Display the first few rows of the dataframe
places.head()
```

	StateAbbr	CountyName	CountyFIPS	LocationName	Category	Measure	D
0	FL	Miami-Dade	12086	12086010025	Health Outcomes	Stroke among adults	
1	FL	Miami-Dade	12086	12086002709	Disability	Vision disability among adults	
2	FL	Miami-Dade	12086	12086007807	Disability	Vision disability among adults	
3	FL	Miami-Dade	12086	12086001206	Health Status	Frequent physical distress among adults	

```
places.info()  
  
places.describe().T
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3183048 entries, 0 to 3183047  
Data columns (total 19 columns):  
 #   Column           Dtype  
 ---  ----  
 0   StateAbbr        object  
 1   CountyName       object  
 2   CountyFIPS       object  
 3   LocationName    int64  
 4   Category         object  
 5   Measure          object  
 6   Data_Value_Unit object  
 7   Data_Value_Type object  
 8   Data_Value       float64  
 9   Low_Confidence_Limit float64  
 10  High_Confidence_Limit float64  
 11  TotalPopulation  int64  
 12  TotalPop18plus  int64  
 13  Geolocation     object  
 14  fips             object  
 15  CategoryID      object  
 16  MeasureId       object  
 17  DataValueTypeID object  
 18  Short_Question_Text object  
dtypes: float64(3), int64(3), object(13)  
memory usage: 461.4+ MB
```

	count	mean	std	min	
LocationName	3183048.0	2.781862e+10	1.601556e+10	1.001020e+09	1.2e+09
Data_Value	3183048.0	2.569415e+01	2.296357e+01	2.000000e-01	9.100000e+00
Low_Confidence_Limit	3183048.0	2.346279e+01	2.192081e+01	2.000000e-01	8.000000e+00
High_Confidence_Limit	3183048.0	2.797654e+01	2.394928e+01	3.000000e-01	1.020000e+01
TotalPopulation	3183048.0	3.972090e+03	1.673262e+03	5.400000e+01	2.770000e+03
TotalPop18plus	3183048.0	3.094050e+03	1.273252e+03	5.000000e+01	2.180000e+03

```
# Ensure 'CountyFIPS' is treated as a string  
places['CountyFIPS'] = places['CountyFIPS'].astype(str)
```

```
# Pivot the DataFrame to have separate columns for DIABETES and STROKE

places_filtered = places[
    places['MeasureId'].isin(["DIABETES", "STROKE"])
][['fips', 'CountyFIPS', 'Measure', 'Data_Value']]
places_pivot = places_filtered.pivot(index=['fips', 'CountyFIPS'], co
places_pivot = places_pivot.reset_index().rename_axis(None, axis=1)

places_pivot.head()
```

	fips	CountyFIPS	Diagnosed diabetes among adults	Stroke among adults
0	01001020100	1001	13.8	4.2
1	01001020200	1001	14.7	4.0
2	01001020300	1001	14.1	4.2
3	01001020400	1001	13.4	4.2
4	01001020501	1001	11.4	3.2

```
# Calculate total Data_Value by CountyFIPS and Measure
total_data_value_by_county_measure = places_filtered.groupby(['CountyFIPS'])
print(total_data_value_by_county_measure.head(10))
```

	CountyFIPS	Measure	Data_Value
0	10001	Diagnosed diabetes among adults	537.3
1	10001	Stroke among adults	168.0
2	10003	Diagnosed diabetes among adults	1878.0
3	10003	Stroke among adults	512.3
4	10005	Diagnosed diabetes among adults	996.6
5	10005	Stroke among adults	323.6
6	1001	Diagnosed diabetes among adults	227.4
7	1001	Stroke among adults	66.5
8	1003	Diagnosed diabetes among adults	589.0
9	1003	Stroke among adults	172.1

```
# Create dataset for 'Diagnosed diabetes among adults'
diabetes_data = total_data_value_by_county_measure[
    total_data_value_by_county_measure['Measure'] == "Diagnosed diabetes among adults"
].copy()
```

```
# Create dataset for 'Stroke among adults'
stroke_data = total_data_value_by_county_measure[
```

```
total_data_value_by_county_measure['Measure'] == "Stroke among adults"]
].copy()

# Prepare places dataset for merging (select relevant columns and drop
county_info = places[['CountyFIPS', 'StateAbbr', 'CountyName']].drop_()

# Merge diabetes_data with county_info
diabetes_data = pd.merge(diabetes_data, county_info, on='CountyFIPS', how='left')

# Merge stroke_data with county_info
stroke_data = pd.merge(stroke_data, county_info, on='CountyFIPS', how='left')

print("Diabetes Data Head:")
print(diabetes_data.head())
print("\nStroke Data Head:")
print(stroke_data.head())
```

Diabetes Data Head:

	CountyFIPS	Measure	Data_Value	StateAbbr	\
0	10001	Diagnosed diabetes among adults	537.3	DE	
1	10003	Diagnosed diabetes among adults	1878.0	DE	
2	10005	Diagnosed diabetes among adults	996.6	DE	
3	1001	Diagnosed diabetes among adults	227.4	AL	
4	1003	Diagnosed diabetes among adults	589.0	AL	

	CountyName
0	Kent
1	New Castle
2	Sussex
3	Autauga
4	Baldwin

Stroke Data Head:

	CountyFIPS	Measure	Data_Value	StateAbbr	CountyName
0	10001	Stroke among adults	168.0	DE	Kent
1	10003	Stroke among adults	512.3	DE	New Castle
2	10005	Stroke among adults	323.6	DE	Sussex
3	1001	Stroke among adults	66.5	AL	Autauga
4	1003	Stroke among adults	172.1	AL	Baldwin

```
# Get unique measures
unique_measures = total_data_value_by_county_measure['Measure'].unique()

plt.figure(figsize=(15, 7))


for i, measure in enumerate(unique_measures):
```

```
# Filter data for the current measure
measure_data = total_data_value_by_county_measure[
    total_data_value_by_county_measure['Measure'] == measure
]

# Get top 3 CountyFIPS for the current measure
top_3_counties = measure_data.nlargest(3, 'Data_Value')

# Create a subplot for each measure
plt.subplot(1, len(unique_measures), i + 1)
sns.barplot(x='CountyFIPS', y='Data_Value', data=top_3_counties)
plt.title(f'Top 3 CountyFIPS for {measure}')
plt.xlabel('CountyFIPS')
plt.ylabel('Total Data Value')
plt.xticks(rotation=45, ha='right')

plt.tight_layout()
plt.show()
```



```
# Function to plot top 5 highest and lowest counties for a given measure

def plot_top_and_bottom_counties(df, measure_name, value_col='Data_Value'):
    # Sort for top 5 highest
    top_5_highest = df.nlargest(5, value_col)
    # Sort for top 5 lowest
    top_5_lowest = df.nsmallest(5, value_col)

    # Create combined label for x-axis
    top_5_highest['County_State'] = top_5_highest['CountyName'] + ' (' + top_5_highest['StateName'] + ')'
    top_5_lowest['County_State'] = top_5_lowest['CountyName'] + ' (' + top_5_lowest['StateName'] + ')'

    fig, axes = plt.subplots(1, 2, figsize=(18, 6))
    fig.suptitle(f'{measure_name} Prevalence: Top 5 Highest and Lowest')

    # Plot top 5 highest
    sns.barplot(
        ax=axes[0],
        x='County_State',
        y=value_col,
        data=top_5_highest,
        hue='County_State', # Assign x to hue
        palette='viridis',
        legend=False # Disable the legend for the same effect
    )
    axes[0].set_title('Top 5 Counties (Highest Prevalence)')
    axes[0].set_xlabel('County (State)')
    axes[0].set_ylabel(f'{measure_name} Data Value')
    axes[0].tick_params(axis='x', rotation=45)

    # Plot top 5 lowest
    sns.barplot(
        ax=axes[1],
```

```


        x='County_State',
        y=value_col,
        data=top_5_lowest,
        hue='County_State', # Assign x to hue
        palette='plasma',
        legend=False # Disable the legend for the same effect
    )
    axes[1].set_title('Top 5 Counties (Lowest Prevalence)')
    axes[1].set_xlabel('County (State)')
    axes[1].set_ylabel(f'{measure_name} Data Value')
    axes[1].tick_params(axis='x', rotation=45)

    plt.tight_layout(rect=[0, 0.03, 1, 0.95])
    plt.show()

# Plot for Diabetes Data
plot_top_and_bottom_counties(diabetes_data, 'Diabetes')

# Plot for Stroke Data
plot_top_and_bottom_counties(stroke_data, 'Stroke')

```




```
acs_vars = [
    "NAME", "B19013_001E", "B17001_002E", "B23025_005E",
    "B15003_022E", "B15003_023E", "B15003_024E", "B15003_025E",
    "B27001_017E"
]
acs_var_string = ",".join(acs_vars)
print(acs_var_string)
# acs_url = f"https://api.census.gov/data/2022/acs/acs5?get=NAME,B19013_001E,B17001_002E,B23025_005E,B15003_022E,B15003_023E,B15003_024E,B15003_025E,B27001_017E"
# download_csv(acs_url, "/Users/sankates/workspace/MSAI/AI-391MCSMLAC/MSAI-AI-391-MCSMLAC-Project/acs5.csv")
```

```
NAME,B19013_001E,B17001_002E,B23025_005E,B15003_022E,B15003_023E,B15003_024E,B15003_025E,B27001_017E
```

```
# Load ACS data
acs = pd.read_json("/Users/sankates/workspace/MSAI/AI-391MCSML/ACS_ra\
acs.columns = acs.iloc[0]
acs = acs[1:]

acs['fips'] = acs['state'] + acs['county']

# Rename columns for clarity
acs = acs.rename(columns={
    "B19013_001E": "median_income",
    "B17001_002E": "poverty_count",
    "B23025_005E": "unemployed_count",
    "B27001_017E": "uninsured_count"
})

acs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3222 entries, 1 to 3222
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   NAME             3222 non-null   object 
 1   median_income    3222 non-null   object 
 2   poverty_count    3222 non-null   object 
 3   unemployed_count 3222 non-null   object 
 4   B15003_022E      3222 non-null   object 
 5   B15003_023E      3222 non-null   object 
 6   B15003_024E      3222 non-null   object 
 7   B15003_025E      3222 non-null   object 
 8   uninsured_count  3222 non-null   object 
 9   state            3222 non-null   object 
 10  county           3222 non-null   object 
 11  fips             3222 non-null   object 
dtypes: object(12)
memory usage: 302.2+ KB
```

```
acs.head()
```

	NAME	median_income	poverty_count	unemployed_count	B15003_022E
1	Autauga County, Alabama	68315	6630	752	6726
2	Baldwin County, Alabama	71039	23445	3825	33474
	Barbour				

```
# ejs_url = "https://gaftp.epa.gov/EJSCREEN/2023/EJSCREEN_2023_USPR_Pi"
# download_csv(ejs_url, "/Users/sankates/workspace/MSAI/AI-391MCSMLEJ"

# with zipfile.ZipFile("EJSCREEN.zip", "r") as z:
#     z.extractall("ejscreen")

bg = pd.read_csv("/Users/sankates/workspace/MSAI/AI-391MCSMLEJ/EJSCREEN.csv")

base_cols = ['ID', 'STATE_NAME', 'ST_ABBREV', 'CNTY_NAME', 'ACSTOTPOP']
missing_base = [c for c in base_cols if c not in bg.columns]
print(missing_base)

print(bg.info())
bg.head(5)
```

```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/308739
bg = pd.read_csv("/Users/sankates/workspace/MSAI/AI-391MCSMLEJ/EJSCREEN.csv")
[1]:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 243022 entries, 0 to 243021
Columns: 229 entries, ID to Shape_Area
dtypes: float64(173), int64(3), object(53)
memory usage: 424.6+ MB
None
```

	ID	STATE_NAME	ST_ABBREV	CNTY_NAME	REGION	ACSTOTPOP	ACSI
0	10010201001	Alabama	AL	Autauga County	4	558.0	
1	10010201002	Alabama	AL	Autauga County	4	1307.0	
2	10010202001	Alabama	AL	Autauga County	4	548.0	
3	10010202002	Alabama	AL	Autauga County	4	1313.0	

```
# Pick candidate feature columns
candidate_features = [
    'PM25', 'OZONE', 'DSLPM', 'RSEI_AIR', 'PTRAF',
    'PNPL', 'PRMP', 'PWDIS', 'DWATER', 'N02',
    'DEMOGIDX_2', 'DEMOGIDX_5',
    'PEOPCOLORPCT', 'LOWINCPCT', 'UNEMPPCT', 'LINGISOPCT', 'LESSHSPCT', 'UN
]

feature_cols = [c for c in candidate_features if c in bg.columns]
print("Using feature columns:", feature_cols)
```

```
Using feature columns: ['PM25', 'OZONE', 'DSLPM', 'RSEI_AIR', 'PTRAF',
```

```
# Aggregate block group data to county level

output_path = "/Users/sankates/workspace/MSAI/AI-391MCSML/EJSCREEN_COI

# Check if the file already exists
if os.path.exists(output_path):
    county_ejs = pd.read_csv(output_path)
    print(county_ejs.info())
    print(f"Loaded county_ejs from existing file: {output_path}")
else:
    # Ensure numeric
    for col in feature_cols + ['ACSTOTPOP']:
        bg[col] = pd.to_numeric(bg[col], errors='coerce')

    # Drop rows without population
    bg = bg[bg['ACSTOTPOP'] > 0].copy()

    group_cols = ['ST_ABBREV', 'STATE_NAME', 'CNTY_NAME']

    def weighted_mean(series, weights):
        mask = series.notna() & weights.notna() & (weights > 0)
        if not mask.any():
            return np.nan
        return np.average(series[mask], weights=weights[mask])

    grouped = bg.groupby(group_cols, as_index=False)

    records = []
    for keys, group in grouped:
        rec = {}
```

```
# keys is tuple matching group_cols
for k, name in zip(keys, group_cols):
    rec[name] = k
rec['total_pop'] = group['ACSTOTPOP'].sum()
for col in feature_cols:
    rec[col] = weighted_mean(group[col], group['ACSTOTPOP'])
records.append(rec)

county_ejs = pd.DataFrame(records)
county_ejs.to_csv(output_path, index=False)
print(f"Generated and saved county_ejs to: {output_path}")

county_ejs.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3233 entries, 0 to 3232
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   ST_ABBREV        3233 non-null   object  
 1   STATE_NAME       3233 non-null   object  
 2   CNTY_NAME        3233 non-null   object  
 3   total_pop        3233 non-null   float64 
 4   PM25             3109 non-null   float64 
 5   OZONE             3109 non-null   float64 
 6   DSLPM             3192 non-null   float64 
 7   RSEI_AIR          3233 non-null   float64 
 8   PTRAF             3222 non-null   float64 
 9   PNPL              3230 non-null   float64 
 10  PRMP              3230 non-null   float64 
 11  PWDIS             3192 non-null   float64 
 12  DWATER            3142 non-null   float64 
 13  NO2               3222 non-null   float64 
 14  DEMOGIDX_2        3233 non-null   float64 
 15  DEMOGIDX_5        3222 non-null   float64 
 16  PEOPCOLORPCT      3233 non-null   float64 
 17  LOWINCPCT         3233 non-null   float64 
 18  UNEMPPCT          3233 non-null   float64 
 19  LINGISOPCT        3233 non-null   float64 
 20  LESSHSPCT          3233 non-null   float64 
 21  UNDER5PCT          3233 non-null   float64 
 22  OVER64PCT          3233 non-null   float64 
dtypes: float64(20), object(3)
memory usage: 581.1+ KB
None
Loaded county_ejs from existing file: /Users/sankates/workspace/MSAI/A1
```

```
ST_ABBREV  STATE_NAME  CNTY_NAME  total_pop  PM25  OZONE  DSLPM  RSEI_
```

0	AK	Alaska	Aleutians East Borough	3407.0	NaN	NaN	NaN	18.79
1	AK	Alaska	Aleutians West Census Area	5219.0	NaN	NaN	NaN	3.63
2	AK	Alaska	Anchorage Municipality	290674.0	NaN	NaN	NaN	47.48

```
# Load FIPS lookup data
```

```
fips_lookup = pd.read_excel('/Users/sankates/workspace/MSAI/AI-391MCSI'
print(fips_lookup.head())
```

```
# Rename columns in fips_lookup for easier merging
```

```
# Corrected column names based on available fips_lookup (from kernel):
fips_lookup = fips_lookup.rename(columns={'Name': 'STATE_NAME', 'fips
```

```
# Create the new dataset by merging county_ejs with fips_lookup
```

```
county_ejs_with_fips = pd.merge(
    county_ejs,
    fips_lookup[['STATE_NAME', 'CNTY_NAME', 'fips_from_lookup']],
    on=['STATE_NAME', 'CNTY_NAME'],
    how='left'
)
```

```
# Rename the merged fips column to 'fips'
```

```
county_ejs_with_fips = county_ejs_with_fips.rename(columns={'fips_fro
```

```
county_ejs_with_fips.tail(10)
```

State	FIPS	Code	County	FIPS	Code	Area	Name	Name	fips
ST_ABBREV	STATE_NAME	CNTY_NAME				total_pop	PM25	OZONE	INCOME
3223	WY	Wyoming		Niobrara County		2460.0	4.460460	64.361420	0.0
3224	WY	Wyoming		Park County		29878.0	4.630895	60.448647	0.0

3225	WY	Wyoming	Platte County	8618.0	4.454283	65.295025	0.0
3226	WY	Wyoming	Sheridan County	31176.0	5.131052	58.859135	0.0
3227	WY	Wyoming	Sublette County	8801.0	3.551251	63.564417	0.0
3228	WY	Wyoming	Sweetwater County	42079.0	4.282689	63.228821	0.1
3229	WY	Wyoming	Teton County	23346.0	4.445018	61.689650	0.0

```
# Convert 'fips' column to int64, handling NaN values
# Fill NaN values with 0 before converting to int, as int type cannot
county_ejs_with_fips['fips'] = county_ejs_with_fips['fips'].fillna(0)

output_csv_path = "/Users/sankates/workspace/MSAI/AI-391MCSML/county_ejs_with_fips.csv"
county_ejs_with_fips.to_csv(output_csv_path, index=False)

print(f"DataFrame saved to {output_csv_path}")
print("Info of saved DataFrame with converted fips column:")
county_ejs_with_fips.info()
county_ejs_with_fips.head()
```

```
DataFrame saved to /Users/sankates/workspace/MSAI/AI-391MCSML/county_ejs_with_fips.csv
Info of saved DataFrame with converted fips column:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3233 entries, 0 to 3232
Data columns (total 24 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ST_ABBREV        3233 non-null   object 
 1   STATE_NAME       3233 non-null   object 
 2   CNTY_NAME        3233 non-null   object 
 3   total_pop        3233 non-null   float64
 4   PM25             3109 non-null   float64
 5   OZONE            3109 non-null   float64
 6   DSLPM            3192 non-null   float64
 7   RSEI_AIR         3233 non-null   float64
 8   PTRAF            3222 non-null   float64
 9   PNPL             3230 non-null   float64
 10  PRMP             3230 non-null   float64
 11  PWDIS            3192 non-null   float64
 12  DWATER           3142 non-null   float64
 ..   ...
```

```

13 NO2            3222 non-null    float64
14 DEMOGIDX_2    3233 non-null    float64
15 DEMOGIDX_5    3222 non-null    float64
16 PEOPCOLORPCT  3233 non-null    float64
17 LOWINCPCT     3233 non-null    float64
18 UNEMPPCT      3233 non-null    float64
19 LINGISOPCT    3233 non-null    float64
20 LESSHSPCT     3233 non-null    float64
21 UNDER5PCT     3233 non-null    float64
22 OVER64PCT     3233 non-null    float64
23 fips          3233 non-null    int64
dtypes: float64(20), int64(1), object(3)
memory usage: 606.3+ KB

```

	ST_ABBREV	STATE_NAME	CNTY_NAME	total_pop	PM25	OZONE	DSLPM	RSEI
0	AK	Alaska	Aleutians East Borough	3407.0	NaN	NaN	NaN	18.79
1	AK	Alaska	Aleutians West Census Area	5219.0	NaN	NaN	NaN	3.63
2	AK	Alaska	Anchorage Municipality	290674.0	NaN	NaN	NaN	47.48

```

# Load Food Access Research Atlas data
food = pd.read_csv("/Users/sankates/workspace/MSAI/AI-391MCSML/Food Access Research Atlas Data.csv")
#print(food.columns)
# CensusTract,State,County,Urban,Pop2010,OHU2010,GroupQuartersFlag,NUU

# Create county FIPS (first 5 digits of tract GEOID)
food['fips'] = food['CensusTract'].str[:5]

# Only use tracts with recorded population
food = food[food['Pop2010'] > 0]

# Functions for weighted averages
def weighted_mean(series, weights):
    mask = series.notna() & weights.notna() & (weights > 0)
    if not mask.any():
        return np.nan
    return np.average(series[mask], weights=weights[mask])

# Aggregate to county level
food_county = (

```

```
food.groupby('fips')
    .apply(lambda g: pd.Series({
        # Count of low-income / low-access tracts
        'low_access_tracts': g['LILATracts_1And10'].fillna(0).sum()

        # Percent of households with low vehicle access
        'pct_low_vehicle_access': weighted_mean(g['HUNVFlag'], g['

        # Poverty / SNAP weighted averages
        'poverty_rate': weighted_mean(g['PovertyRate'], g['Pop2010
    }), include_groups=False)
    .reset_index()
)

food_county.head()
```

	fips	low_access_tracts	pct_low_vehicle_access	poverty_rate
0	10001	7.0	0.526215	13.881833
1	10003	19.0	0.239005	11.927730
2	10005	6.0	0.070628	11.873381
3	10010	5.0	0.449341	15.137584
4	10030	4.0	0.223543	10.947872

```
# Ensure 'fips' is numeric
food_county['fips'] = pd.to_numeric(food_county['fips'], errors='coerce')
print(food_county.nsmallest(5, 'fips'))
```

	fips	low_access_tracts	pct_low_vehicle_access	poverty_rate
0	10001	7.0	0.526215	13.881833
1	10003	19.0	0.239005	11.927730
2	10005	6.0	0.070628	11.873381
3	10010	5.0	0.449341	15.137584
4	10030	4.0	0.223543	10.947872

```
# Load ERS County Typology data
df = pd.read_csv("/Users/sankates/workspace/MSAI/AI-391MCSML/erscounty.csv")

# Pivot Attribute → feature columns
rucc = df.pivot_table(
    index=["FIPStxt", "State", "County_Name", "Metro2023"],
    columns="Attribute",
    values="Value",
    aggfunc="first"
).reset_index()

# Clean up column names
rucc.columns.name = None
rucc = rucc.rename(columns={"FIPStxt": "fips"})

# Show result
rucc.head()
```

	fips	State	County_Name	Metro2023	High_Farming_2025	High_Governm...
0	01001	AL	Autauga County	1	0	
1	01003	AL	Baldwin County	1	0	
2	01005	AL	Barbour County	0	0	
3	01007	AL	Bibb County	1	0	
4	01009	AL	Blount County	1	0	

```
from functools import reduce

# --- 1. Standardize 'fips' column and select relevant columns for each county

# Prepare rucc
rucc_standardized = rucc.copy()
# Ensure fips is numeric and int64, handle potential non-numeric values
rucc_standardized['fips'] = pd.to_numeric(rucc_standardized['fips'], errors='coerce')
# Drop redundant state/county name columns as they will come from county typology
```

```
rucc_standardized = rucc_standardized.drop(columns=['State', 'County_I']

# Prepare food_county (fips is already int64 after previous steps)
food_county_standardized = food_county.copy()
food_county_standardized['fips'] = food_county_standardized['fips'].ast

# Prepare county_ejs_with_fips (fips is already int64 after previous steps)
# This will serve as our base DataFrame for merging.
county_ejs_standardized = county_ejs_with_fips.copy()
county_ejs_standardized['fips'] = county_ejs_standardized['fips'].ast

# Prepare acs
acs_standardized = acs[['fips', 'median_income', 'poverty_count', 'unem
acs_standardized['fips'] = pd.to_numeric(acs_standardized['fips'], errors='

# Prepare diabetes_data
diabetes_standardized = diabetes_data.copy()
diabetes_standardized = diabetes_standardized.rename(columns={'CountyI
# Ensure fips is int64, handle potential non-numeric values by filling
diabetes_standardized['fips'] = pd.to_numeric(diabetes_standardized['fips
# Select only fips and the new prevalence column
diabetes_standardized = diabetes_standardized[['fips', 'diabetes_prev

# Prepare stroke_data
stroke_standardized = stroke_data.copy()
stroke_standardized = stroke_standardized.rename(columns={'CountyFIPS
# Ensure fips is int64, handle potential non-numeric values by filling
stroke_standardized['fips'] = pd.to_numeric(stroke_standardized['fips
# Select only fips and the new prevalence column
stroke_standardized = stroke_standardized[['fips', 'stroke_prev']]

# --- 2. List of DataFrames to merge, starting with county_ejs_with_fips
dfs_to_merge = [
    county_ejs_standardized, # Base dataframe
    rucc_standardized,
    food_county_standardized,
    acs_standardized,
    diabetes_standardized,
    stroke_standardized
]

# --- 3. Perform the merge ---
data = reduce(lambda left, right: pd.merge(left, right, on='fips', how='left', sort=False), dfs_to_merge)

# Display head to confirm the merge result
```

```
print("Merged Data Head:")
data.head()
```

Merged Data Head:

	ST_ABBREV	STATE_NAME	CNTY_NAME	total_pop	PM25	OZONE	DSLPM	RSEI_AIR
0	AK	Alaska	Aleutians East Borough	3407.0	NaN	NaN	NaN	18.79
1	AK	Alaska	Aleutians West Census Area	5219.0	NaN	NaN	NaN	3.63
2	AK	Alaska	Anchorage Municipality	290674.0	NaN	NaN	NaN	47.48
3	AK	Alaska	Bethel Census Area	18538.0	NaN	NaN	NaN	24.38

```
output_path = "/Users/sankates/workspace/MSAI/AI-391MCSML/data.csv"
data.to_csv(output_path, index=False)
print(f"DataFrame 'data' saved to {output_path}")
```

DataFrame 'data' saved to /Users/sankates/workspace/MSAI/AI-391MCSML/d

```
# Display info of the final merged DataFrame
print("Merged Data Info:")
data.info()
```

Merged Data Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 3233 entries, 0 to 3232

Data columns (total 47 columns):

#	Column	Non-Null Count	Dtype
0	ST_ABBREV	3233 non-null	object
1	STATE_NAME	3233 non-null	object
2	CNTY_NAME	3233 non-null	object
3	total_pop	3233 non-null	float64
4	PM25	3109 non-null	float64
5	OZONE	3109 non-null	float64
6	DSLPM	3192 non-null	float64
7	RSEI_AIR	3233 non-null	float64

```
8    PTRAF                3222 non-null  float64
9    PNPL                 3230 non-null  float64
10   PRMP                 3230 non-null  float64
11   PWDIS                3192 non-null  float64
12   DWATER                3142 non-null  float64
13   N02                  3222 non-null  float64
14   DEMOGIDX_2            3233 non-null  float64
15   DEMOGIDX_5            3222 non-null  float64
16   PEOPCOLORPCT          3233 non-null  float64
17   LOWINCPCT              3233 non-null  float64
18   UNEMPPCT              3233 non-null  float64
19   LINGISOPCT              3233 non-null  float64
20   LESSHSPCT              3233 non-null  float64
21   UNDER5PCT              3233 non-null  float64
22   OVER64PCT              3233 non-null  float64
23   fips                  3233 non-null  int64
24   Metro2023              3095 non-null  float64
25   High_Farming_2025        3095 non-null  float64
26   High_Government_2025        3095 non-null  float64
27   High_Manufacturing_2025        3095 non-null  float64
28   High_Mining_2025            3095 non-null  float64
29   High_Recreation_2025          3095 non-null  float64
30   Housing_Stress_2025          3095 non-null  float64
31   Industry_Dependence_2025        3095 non-null  float64
32   Low_Employment_2025            3095 non-null  float64
33   Low_PostSecondary_Ed_2025        3095 non-null  float64
34   Nonspecialized_2025            3095 non-null  float64
35   Persistent_Poverty_1721          3095 non-null  float64
36   Population_Loss_2025            3095 non-null  float64
37   Retirement_Destination_2025        3095 non-null  float64
38   low_access_tracts            2786 non-null  float64
39   pct_low_vehicle_access          2786 non-null  float64
40   poverty_rate                2785 non-null  float64
41   median_income                3095 non-null  object
42   poverty_count                3095 non-null  object
43   unemployed_count              3095 non-null  object
44   uninsured_count              3095 non-null  object
45   diabetes_prev                3094 non-null  float64
46   stroke_prev                  3094 non-null  float64
dtypes: float64(39), int64(1), object(7)
memory usage: 1.2+ MB
```

```
data.head()
```

	ST_ABBREV	STATE_NAME	CNTY_NAME	total_pop	PM25	OZONE	DSLPM	RSEI_
0	AK	Alaska	Aleutians East Borough	3407.0	NaN	NaN	NaN	18.79
1	AK	Alaska	Aleutians West Census Area	5219.0	NaN	NaN	NaN	3.63
2	AK	Alaska	Anchorage Municipality	290674.0	NaN	NaN	NaN	47.48
3	AK	Alaska	Bethel Census Area	18538.0	NaN	NaN	NaN	24.38

```
data.shape
```

```
(3233, 47)
```

```
# Create a copy of the original 'data' DataFrame
data_cleaned = data.copy()

# 1. Change specified columns to int64
columns_to_convert_to_int = ['median_income', 'poverty_count', 'unemp']

for col in columns_to_convert_to_int:
    # Convert to numeric, coercing errors to NaN
    data_cleaned[col] = pd.to_numeric(data_cleaned[col], errors='coerce')
    # Fill NaN values with 0 before converting to int64
    data_cleaned[col] = data_cleaned[col].fillna(0).astype('int64')

# 2. Drop specified columns
columns_to_drop = ['CNTY_NAME', 'STATE_NAME', 'ST_ABBREV']
data_cleaned = data_cleaned.drop(columns=columns_to_drop)

print("Info of the new data_cleaned DataFrame:")
data_cleaned.info()
print("\nHead of the new data_cleaned DataFrame:")
data_cleaned.head()
```

Info of the new data_cleaned DataFrame:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 3233 entries, 0 to 3232

Data columns (total 44 columns):

#	Column	Non-Null Count	Dtype
0	total_pop	3233	float64
1	PM25	3109	float64
2	OZONE	3109	float64
3	DSLPM	3192	float64
4	RSEI_AIR	3233	float64
5	PTRAF	3222	float64
6	PNPL	3230	float64
7	PRMP	3230	float64
8	PWDIS	3192	float64
9	DWATER	3142	float64
10	NO2	3222	float64
11	DEMOGIDX_2	3233	float64
12	DEMOGIDX_5	3222	float64
13	PEOPCOLORPCT	3233	float64
14	LOWINCPECT	3233	float64
15	UNEMPPCT	3233	float64
16	LINGISOPCT	3233	float64
17	LESSHSPCT	3233	float64
18	UNDER5PCT	3233	float64
19	OVER64PCT	3233	float64
20	fips	3233	non-null
21	Metro2023	3095	float64
22	High_Farming_2025	3095	float64
23	High_Government_2025	3095	float64
24	High_Manufacturing_2025	3095	float64
25	High_Mining_2025	3095	float64
26	High_Recreation_2025	3095	float64
27	Housing_Stress_2025	3095	float64
28	Industry_Dependence_2025	3095	float64
29	Low_Employment_2025	3095	float64
30	Low_PostSecondary_Ed_2025	3095	float64
31	Nonspecialized_2025	3095	float64
32	Persistent_Poverty_1721	3095	float64
33	Population_Loss_2025	3095	float64
34	Retirement_Destination_2025	3095	float64
35	low_access_tracts	2786	float64
36	pct_low_vehicle_access	2786	float64
37	poverty_rate	2785	float64
38	median_income	3233	non-null
39	poverty_count	3233	non-null
40	unemployed_count	3233	non-null
41	uninsured_count	3233	non-null
42	diabetes_prev	3094	float64
43	stroke_prev	3094	float64

```
dtypes: float64(39), int64(5)
memory usage: 1.1 MB
```

Head of the new data_cleaned DataFrame:

	total_pop	PM25	OZONE	DSLPM	RSEI_AIR	PTRAF	PNPL	PRI
0	3407.0	NaN	NaN	NaN	18.795745	0.000000e+00	0.000000	2.3570
1	5219.0	NaN	NaN	NaN	3.631383	0.000000e+00	0.323449	0.1881
2	290674.0	NaN	NaN	NaN	47.484369	1.186345e+06	1.337817	0.4477
3	18538.0	NaN	NaN	NaN	24.384799	0.000000e+00	0.000000	0.0000
4	854.0	NaN	NaN	NaN	22.729112	0.000000e+00	0.000000	0.1925

5 rows x 44 columns

```
# Load merged dataset
df = data_cleaned.copy()

# Select only numeric columns for correlation
numeric_df = df.select_dtypes(include=[np.number])


numeric_df["poverty_rate"] = numeric_df["poverty_count"] / numeric_df["total_pop"]
numeric_df["unemployment_rate"] = numeric_df["unemployed_count"] / numeric_df["total_pop"]
numeric_df["uninsured_rate"] = numeric_df["uninsured_count"] / numeric_df["total_pop"]
numeric_df["total_pop"] = np.log1p(df["total_pop"])

# Drop poverty_count, unemployed_count, uninsured_count
numeric_df = numeric_df.drop(columns=["poverty_count", "unemployed_count", "uninsured_count"])

# Compute correlation matrix
corr = numeric_df.corr()

# Plot heatmap (matplotlib only, no specified colors)
plt.figure(figsize=(12, 10))
plt.imshow(corr, aspect='auto')
plt.colorbar()
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.columns)), corr.columns)
plt.title("Correlation Heatmap of Independent Variables")

plt.tight_layout()
plt.show()
```



```
# Targets
targets = ["diabetes_prev", "stroke_prev"]

ranked_results = {}

for target in targets:
    if target in corr.columns:
        ranked = corr[target].drop(labels=[target]).abs().sort_values()
        ranked_results[target] = ranked

# Extract top 20 predictors for diabetes_prev
top_diabetes = ranked_results["diabetes_prev"].head(20)
top_diabetes
```

stroke_prev	0.994584
low_access_tracts	0.752232
PTRAF	0.586347
DSLPM	0.557147
total_pop	0.530519
N02	0.405743
Metro2023	0.296897
LINGISOPCT	0.294815
Housing_Stress_2025	0.290966
OZONE	0.283500
PEOPCOLORPCT	0.259983
PM25	0.221053
median_income	0.210882
PRMP	0.209355
OVER64PCT	0.183075
Nonspecialized_2025	0.169826
unemployment_rate	0.153201
High_Farming_2025	0.120520
Industry_Dependence_2025	0.117146
DEMOGIDX_2	0.108128
Name: diabetes_prev, dtype:	float64

```
# Extract top 20 predictors for stroke_prev
top_stroke = ranked_results["stroke_prev"].head(20)
top_stroke
```

```
diabetes_prev          0.994584
low_access_tracts      0.737570
PTRAF                 0.604597
DSLPM                 0.577473
total_pop              0.554268
N02                   0.415541
Metro2023              0.310801
Housing_Stress_2025    0.303543
LINGISOPCT              0.286679
OZONE                  0.285590
PEOPCOLORPCT             0.257388
PM25                   0.223269
median_income             0.213146
PRMP                   0.207583
OVER64PCT                0.178539
Nonspecialized_2025      0.173761
unemployment_rate        0.163062
High_Farming_2025        0.129321
Industry_Dependence_2025  0.117552
DEMOGIDX_2                0.105260
Name: stroke_prev, dtype: float64
```

```
# Load merged dataset
df = data_cleaned.copy()

# Columns that represent raw counts
count_cols = [
    "poverty_count", "unemployed_count", "uninsured_count",
    "stroke_prev", "diabetes_prev", 'High_Manufacturing_2025', 'Low_Er
]

# Check existence
existing = [c for c in count_cols if c in df.columns]

# Create log-transformed population
if "total_pop" in df.columns:
    df["log_total_pop"] = np.log1p(df["total_pop"])
else:
    raise ValueError("total_pop column missing.")

# Create rates per 100k
for col in existing:
    df[col + "_rate"] = df[col] / df["total_pop"] * 100000

# Drop raw count columns safely
df_clean = df.drop(columns=existing)

# Save cleaned file
output_path = "/Users/sankates/workspace/MSAI/AI-391MCSML/data_normal:
df_clean.to_csv(output_path, index=False)
```

```
df_clean.head()
```

	total_pop	PM25	OZONE	DSLPM	RSEI_AIR	PTRAF	PNPL	PRMP
0	3407.0	NaN	NaN	NaN	18.795745	0.000000e+00	0.000000	2.357087
1	5219.0	NaN	NaN	NaN	3.631383	0.000000e+00	0.323449	0.188127
2	290674.0	NaN	NaN	NaN	47.484369	1.186345e+06	1.337817	0.447789
3	18538.0	NaN	NaN	NaN	24.384799	0.000000e+00	0.000000	0.000000
4	854.0	NaN	NaN	NaN	22.729112	0.000000e+00	0.000000	0.192515

5 rows × 45 columns


```
# Create heatmap of correlations in normalized dataset


# Select numeric columns only
numeric_df = df_clean.select_dtypes(include=[np.number])

# Drop specified column
if 'Nonspecialized_2025' in numeric_df.columns:
    numeric_df = numeric_df.drop(columns=['Nonspecialized_2025'])

# Compute correlation matrix
corr = numeric_df.corr()

# Plot heatmap
plt.figure(figsize=(16, 14))
plt.imshow(corr, aspect='auto')
plt.colorbar()
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.columns)), corr.columns)
plt.title("Correlation Heatmap (Normalized Dataset)")
plt.tight_layout()
plt.show()
```





```
# Modeling script

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, ElasticNet
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRe
from sklearn.metrics import mean_squared_error, r2_score, root_mean_sq
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score

# -----
# 0. Load your merged dataset
# -----
df = df_clean.copy()

# targets (dependent variables)
y_diab = df["diabetes_prev_rate"]
y_stroke = df["stroke_prev_rate"]
```

```
# drop rows where either target is missing
mask = y_diab.notna() & y_stroke.notna()
df = df[mask].copy()
y_diab = y_diab[mask]
y_stroke = y_stroke[mask]

# Winsorize (cap at 99th percentile) as a quick fix
from scipy.stats import mstats
y_diab_winsor = mstats.winsorize(y_diab, limits=[0, 0.05]) # Cap top
y_stroke_winsor = mstats.winsorize(y_stroke, limits=[0, 0.05])

# -----
# 2. Define feature matrix X
# -----
drop_cols = set(["diabetes_prev_rate",
                  "stroke_prev_rate", "total_pop", "log_t",
                  "fips"]) # fips is an ID, not a feature

numeric_cols = df.select_dtypes(include=[np.number]).columns
feature_cols = [c for c in numeric_cols if c not in drop_cols]

X = df[feature_cols]

# simple mean imputation for any remaining NaNs
X = X.fillna(X.mean())

# -----
# 3. Train-test split (shared for both targets)
# -----
X_train, X_test, yD_train, yD_test, yS_train, yS_test = train_test_sp
    X, y_diab_winsor, y_stroke_winsor, test_size=0.2, random_state=42
)

# -----
# 4. Helper to train + evaluate a model
# -----
def eval_model(model, X_tr, X_te, y_tr, y_te):
    model.fit(X_tr, y_tr)
    preds = model.predict(X_te)
    rmse = np.sqrt(mean_squared_error(y_te, preds))
    r2 = r2_score(y_te, preds)
    return rmse, r2
```

```
results = []

# -----
# 5. MODELS
# -----


# 5.1 Linear Regression (needs scaling)
lin_pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("linreg", LinearRegression())
])

results.append(["LinearRegression", "diabetes"] +
    list(eval_model(lin_pipeline, X_train, X_test, yD_train, yD_test)))
results.append(["LinearRegression", "stroke"] +
    list(eval_model(lin_pipeline, X_train, X_test, yS_train, yS_test)))

# 5.2 Random Forest (nonlinear, handles interactions)
rf = RandomForestRegressor(
    n_estimators=200,
    random_state=42,
    n_jobs=-1
)

results.append(["RandomForest", "diabetes"] +
    list(eval_model(rf, X_train, X_test, yD_train, yD_test)))
results.append(["RandomForest", "stroke"] +
    list(eval_model(rf, X_train, X_test, yS_train, yS_test)))

# 5.3 Gradient Boosting (you can swap in XGBoost if installed)
#gb = GradientBoostingRegressor(random_state=42)
gb = XGBRegressor(n_estimators=300, learning_rate=0.05, random_state=42)

results.append(["XGBoost", "diabetes"] +
    list(eval_model(gb, X_train, X_test, yD_train, yD_test)))
results.append(["XGBoost", "stroke"] +
    list(eval_model(gb, X_train, X_test, yS_train, yS_test)))

# 5.4 Elastic Net (L1 + L2 regularized linear model)
en_pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("enet", ElasticNet(alpha=0.1, l1_ratio=0.5,
        random_state=42, max_iter=10000))
])



```

```
results.append(["ElasticNet", "diabetes"] +  
    list(eval_model(en_pipeline, X_train, X_test, yD_train))  
results.append(["ElasticNet", "stroke"] +  
    list(eval_model(en_pipeline, X_train, X_test, yS_train))  
  
# -----  
# 6. Compare performance  
# -----  
results_df = pd.DataFrame(results,  
    columns=["Model", "Target", "RMSE", "R2"])  
  
print(results_df.sort_values(["Target", "R2"], ascending=[True, False]))
```

	Model	Target	RMSE	R2
2	RandomForest	diabetes	97.033926	0.655837
4	XGBoost	diabetes	97.090497	0.655435
6	ElasticNet	diabetes	103.114787	0.611349
0	LinearRegression	diabetes	103.861457	0.605700
5	XGBoost	stroke	30.310563	0.704462
3	RandomForest	stroke	30.932333	0.692212
1	LinearRegression	stroke	32.652272	0.657033
7	ElasticNet	stroke	32.722188	0.655563

```
# Plotting model performance  
  
# Make sure results_df exists  
results_df = pd.DataFrame(results, columns=["Model", "Target", "RMSE", "R2"])  
results_df = results_df.sort_values(["Target", "R2"], ascending=[True, False])  
  
# Set style  
sns.set_style("whitegrid")  
fig, axes = plt.subplots(1, 2, figsize=(14, 6), sharey=False)  
  
# R2 plot  
sns.barplot(data=results_df, x="R2", y="Model", hue="Target", ax=axes[0])  
axes[0].set_title("R2 Score by Model and Target", fontsize=14, pad=20)  
axes[0].set_xlabel("R2 (higher is better)")  
axes[0].legend(title="Target")  
  
# RMSE plot  
sns.barplot(data=results_df, x="RMSE", y="Model", hue="Target", ax=axes[1])  
axes[1].set_title("RMSE by Model and Target", fontsize=14, pad=20)  
axes[1].set_xlabel("RMSE (lower is better)")  
axes[1].legend(title="Target")
```

```
    plt.tight_layout()
    plt.show()
```



▼ Hyperparameter tuning for the top 2 models

```
# Random Forest hyperparameters
rf_param_dist = {
    "n_estimators": [100, 200, 400],
```

```
"max_depth": [10, 15, 20, None],
"max_features": ["sqrt", "log2", 0.3, 0.5, 0.7],
"min_samples_split": [2, 5, 10],
"min_samples_leaf": [1, 2, 4],
"bootstrap": [True, False],
}

# XGBoost Boosting hyperparameters
xgb_param_dist = {
    "n_estimators": [200, 400, 600, 1000], # Higher for early stopping
    "learning_rate": [0.01, 0.03, 0.05, 0.1],
    "max_depth": [2, 3, 4, 5],
    "min_child_weight": [1, 3, 5], # Replace min_samples_leaf
    "gamma": [0, 0.1, 0.2], # For split regularization (like min_samples_leaf)
    "subsample": [0.6, 0.8, 1.0],
    "colsample_bytree": [0.6, 0.8, 1.0], # Add this
}
from sklearn.model_selection import RandomizedSearchCV

def tune_and_eval(name, estimator, param_dist, X_train, y_train, X_test,
                  n_iter=25):
    search = RandomizedSearchCV(
        estimator,
        param_distributions=param_dist,
        n_iter=n_iter,
        scoring="r2",
        cv=3,
        random_state=42,
        n_jobs=-1,
        verbose=1,
    )
    search.fit(X_train, y_train)
    best_model = search.best_estimator_
    preds = best_model.predict(X_test)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
    r2 = r2_score(y_test, preds)
    return {
        "name": name,
        "best_params": search.best_params_,
        "rmse": rmse,
        "r2": r2,
        "model": best_model,
    }
```

```
# Example: Tune Random Forest for diabetes_prev_rate
results = []

# ---- Diabetes models ----
rf_diab = tune_and_eval(
    "RandomForest_diabetes",
    RandomForestRegressor(random_state=42),
    rf_param_dist,
    X_train, yD_train, X_test, yD_test,
    n_iter=25
)
results.append(rf_diab)

gb_diab = tune_and_eval(
    "XGBoost_diabetes",
    XGBRegressor(random_state=42),
    xgb_param_dist,
    X_train, yD_train, X_test, yD_test,
    n_iter=25
)
results.append(gb_diab)

# ---- Stroke models ----
rf_stroke = tune_and_eval(
    "RandomForest_stroke",
    RandomForestRegressor(random_state=42),
    rf_param_dist,
    X_train, yS_train, X_test, yS_test,
    n_iter=25
)
results.append(rf_stroke)

gb_stroke = tune_and_eval(
    "XGBoost_stroke",
    XGBRegressor(random_state=42),
    xgb_param_dist,
    X_train, yS_train, X_test, yS_test,
    n_iter=25
)
results.append(gb_stroke)

# Display tuning results
for res in results:
    print(f"Model: {res['name']}")
```

```
print("Best Parameters:")
for param, value in res['best_params'].items():
    print(f" - {param}: {value}")
print(f"R²: {res['r2']:.4f}, RMSE: {res['rmse']:.4f}")
print("-" * 40)
```

Fitting 3 folds for each of 25 candidates, totalling 75 fits
Fitting 3 folds for each of 25 candidates, totalling 75 fits
Fitting 3 folds for each of 25 candidates, totalling 75 fits
Fitting 3 folds for each of 25 candidates, totalling 75 fits

Model: RandomForest_diabetes

Best Parameters:

- n_estimators: 400
- min_samples_split: 2
- min_samples_leaf: 1
- max_features: 0.3
- max_depth: 15
- bootstrap: False

R²: 0.6646, RMSE: 95.7949

Model: XGBoost_diabetes

Best Parameters:

- subsample: 0.6
- n_estimators: 1000
- min_child_weight: 3
- max_depth: 5
- learning_rate: 0.01
- gamma: 0.1
- colsample_bytree: 0.8

R²: 0.6800, RMSE: 93.5707

Model: RandomForest_stroke

Best Parameters:

- n_estimators: 400
- min_samples_split: 2
- min_samples_leaf: 1
- max_features: 0.3
- max_depth: 15
- bootstrap: False

R²: 0.7036, RMSE: 30.3569

Model: XGBoost_stroke

Best Parameters:

- subsample: 0.6
- n_estimators: 1000
- min_child_weight: 3
- max_depth: 5
- learning_rate: 0.01

```
- gamma: 0.1
- colsample_bytree: 0.8
R2: 0.7176, RMSE: 29.6300
```


```
# Plot feature importances for the best models


for res in results:
    best_model = res['model']
    model_name = res['name']

    # Get feature importances as a sorted Series
    importances = pd.Series(
        best_model.feature_importances_,
        index=X.columns
    ).sort_values(ascending=False) # Sort descending for top features

    # Plot
    plt.figure(figsize=(12, 8)) # Adjust size for readability
    importances.plot(kind='barh', color='skyblue')
    plt.title(f"Feature Importances for {model_name}")
    plt.xlabel("Importance Score")
    plt.ylabel("Features")
    plt.gca().invert_yaxis() # Highest importance at the top
    plt.tight_layout()
    plt.show()

    # Optional: Print top 10 for quick view
    print(f"Top 10 Features for {model_name}:")
    print(importances.head(10))
    print("\n" + "-" * 50 + "\n")
```






Top 10 Features for RandomForest_diabetes:

PTRAF	0.174937
median_income	0.138520
Population_Loss_2025	0.084540
DSLPM	0.066507
OVER64PCT	0.053423
LOWINCPCT	0.051559
poverty_count_rate	0.040006
NO2	0.035530
PEOPCOLORPCT	0.026061
LESSHSPCT	0.024879


dtype: float64



Top 10 Features for XGBoost_diabetes:

Population_Loss_2025	0.160859
PTRAF	0.139243
median_income	0.104816


```
-----  
High_Farming_2025           0.057504  
OVER64PCT                   0.030062  
poverty_count_rate          0.025773  
PEOPCOLORPCT                0.025360  
pct_low_vehicle_access      0.023241  
DSLPM                        0.022884  
NO2                          0.022475  
dtype: float32  
-----
```



Top 10 Features for RandomForest_stroke:

```
PTRAF                      0.184603  
median_income                0.126162  
DSLPM                       0.104185  
OVER64PCT                   0.080139  
Population_Loss_2025         0.070341  
LOWINCPCT                   0.045969  
NO2                          0.037371  
poverty_count_rate           0.034504  
DEMOGIDX_5                  0.023977  
Metro2023                    0.020745  
dtype: float64  
-----
```

Feature Importances for XGBoost_stroke



Top 10 Features for XGBoost_stroke:

```

Population_Loss_2025      0.155176
PTRAF                      0.150960
median_income                0.097816
High_Farming_2025            0.047284
Metro2023                    0.042388
OVER64PCT                     0.036477
DSLPM                        0.032674
NO2                           0.023365
poverty_count_rate            0.022645
DEMOGIDX_2                     0.022352
dtype: float32

```



```
# Summarize tuning results

summary = pd.DataFrame([
    {
        "Model": r["name"],
        "Target": "diabetes" if "diabetes" in r["name"] else "stroke",
        "RMSE": r["rmse"],
        "R2": r["r2"]
    }
    for r in results
])

print(summary.sort_values(["Target", "R2"], ascending=[True, False]))
```

	Model	Target	RMSE	R2
1	XGBoost_diabetes	diabetes	93.570742	0.679965
0	RandomForest_diabetes	diabetes	95.794908	0.664570
3	XGBoost_stroke	stroke	29.629978	0.717585
2	RandomForest_stroke	stroke	30.356879	0.703558

```
from sklearn.ensemble import VotingRegressor

# Diabetes example (repeat for stroke with yS)
ensemble_diab = VotingRegressor(estimators=[
    ('xgb', results[1]['model']), # Assuming results list order from
    ('rf', results[0]['model'])
])
ensemble_diab.fit(X_train, yD_train) # Or full X/y for final model
preds_ens = ensemble_diab.predict(X_test)
rmse_ens = root_mean_squared_error(yD_test, preds_ens)
r2_ens = r2_score(yD_test, preds_ens)
print(f"Ensemble Diabetes: RMSE={rmse_ens:.2f}, R²={r2_ens:.4f}")

Ensemble Diabetes: RMSE=93.76, R²=0.6787
```

```
import shap
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

# Function to compute and plot SHAP for a model
def shap_analysis(model, model_name, X_data, top_features_only=False):
    # Use TreeExplainer for efficiency on tree models (RF or XGB)
```

```
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_data)

# Summary plot: Shows global feature importance with direction (beeswarm)
plt.figure()
shap.summary_plot(shap_values, X_data, plot_type="dot", show=False)
plt.title(f"SHAP Summary (Beeswarm) for {model_name}")
plt.show()

# Bar plot: Mean absolute SHAP values (global importance with magnitude)
plt.figure()
shap.summary_plot(shap_values, X_data, plot_type="bar", show=False)
plt.title(f"SHAP Mean Absolute Importance for {model_name}")
plt.show()

# If top_features_only=True, filter to top 10 from your lists
if top_features_only:
    # Example: Limit to top 10 features (replace with your actual top 10)
    top_indices = np.argsort(-np.abs(shap_values).mean(0))[:10]
    shap_values_top = shap_values[:, top_indices]
    X_top = X_data.iloc[:, top_indices] if isinstance(X_data, pd.DataFrame) else X_data
    shap.summary_plot(shap_values_top, X_top, show=False)
    plt.title(f"SHAP for Top 10 Features in {model_name}")
    plt.show()

# Return SHAP values for further analysis (e.g., per county)
return shap_values, explainer.expected_value # Base value for prediction

# Run for each model (use X_test or a random subset for speed; e.g., X_train)
# Ensure X_test has column names for better plots
X_shap = X_test.copy() # Or X for global view, but test set is better

# Diabetes models
rf_diab_shap, rf_diab_base = shap_analysis(results[0]['model'], "Random Forest")
xgb_diab_shap, xgb_diab_base = shap_analysis(results[1]['model'], "XGBoost")

# Stroke models
rf_stroke_shap, rf_stroke_base = shap_analysis(results[2]['model'], "Random Forest")
xgb_stroke_shap, xgb_stroke_base = shap_analysis(results[3]['model'], "XGBoost")


# Example: Dependence plot for a key feature (e.g., PTRAF's impact)
shap.dependence_plot("PTRAF", xgb_diab_shap, X_shap, show=False) # Show dependence plot
plt.title("SHAP Dependence for PTRAF in XGBoost_diabetes")
plt.show()
```


```
# Set matplotlib style to a white background theme (e.g., 'default' or
# plt.style.use('default') # 'default' typically gives white background;

# Optional: Local explanation for a specific county (row index in X_shap)
# Get index of a county to explain based on FIPS
idx = data.index[data['fips'] == 6037] # Example FIPS for Los Angeles County
instance_idx = idx[0] if len(idx) > 0 else 0 # Use found index or default
shap.initjs() # For interactive force plots
force_plot = shap.force_plot(xgb_diab_base, xgb_diab_shap[instance_idx, :])
display(force_plot) # In Jupyter, this shows an interactive plot
```


```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/387051.ipynb
```



```
shap.summary_plot(shap_values, X_data, plot_type="dot", show=False)
```







```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/387051
shap.summary_plot(shap_values, X_data, plot_type="bar", show=False)
```









```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/387051
shap.summary_plot(shap_values, x_data, plot_type="dot", show=False)
```







```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/387051
shap.summary_plot(shap_values, X_data, plot_type="bar", show=False)
```


SHAP Mean Absolute Importance for RandomForest_stroke





```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/387051
shap.summary_plot(shap_values, X_data, plot_type="dot", show=False)
```





```
/var/folders/xv/jppz3tnj7z10dvcn0jp_xfk00000gp/T/ipykernel_60556/387051
shap.summary_plot(shap_values, X_data, plot_type="bar", show=False)
```

